

Divide and conquer algorithms

General template

There are many ways to design algorithms.

For example, insertion sort is *incremental*: having sorted $A[1 \dots j-1]$, place $A[j]$ correctly, so that $A[1 \dots j]$ is sorted.

Divide and conquer

Another common approach.

Divide the problem into a number of subproblems.

Conquer the subproblems by solving them recursively.

Base case: If the subproblems are small enough, just solve them by brute force.

Combine the subproblem solutions to give a solution to the original problem.

To sort $A[p \dots r]$:

Divide by splitting into two subarrays $A[p \dots q]$ and $A[q + 1 \dots r]$, where q is the halfway point of $A[p \dots r]$.

Conquer by recursively sorting the two subarrays $A[p \dots q]$ and $A[q + 1 \dots r]$.

Combine by merging the two sorted subarrays $A[p \dots q]$ and $A[q + 1 \dots r]$ to produce a single sorted subarray $A[p \dots r]$. To accomplish this step, we'll define a procedure $\text{MERGE}(A, p, q, r)$.

The recursion bottoms out when the subarray has just 1 element, so that it's trivially sorted.

Merge sort

MERGE-SORT(A, p, r)

if $p < r$

▷ Check for base case

then $q \leftarrow \lfloor (p + r)/2 \rfloor$

▷ Divide

MERGE-SORT(A, p, q)

▷ Conquer

MERGE-SORT($A, q + 1, r$)

▷ Conquer

MERGE(A, p, q, r)

▷ Combine


Initial call: MERGE-SORT($A, 1, n$)

```

MERGE( $A, p, q, r$ )
 $n_1 \leftarrow q - p + 1$ 
 $n_2 \leftarrow r - q$ 
create arrays  $L[1 \dots n_1 + 1]$  and  $R[1 \dots n_2 + 1]$ 
for  $i \leftarrow 1$  to  $n_1$ 
    do  $L[i] \leftarrow A[p + i - 1]$ 
for  $j \leftarrow 1$  to  $n_2$ 
    do  $R[j] \leftarrow A[q + j]$ 
 $L[n_1 + 1] \leftarrow \infty$ 
 $R[n_2 + 1] \leftarrow \infty$ 
 $i \leftarrow 1$ 
 $j \leftarrow 1$ 
for  $k \leftarrow p$  to  $r$ 
    do if  $L[i] \leq R[j]$ 
        then  $A[k] \leftarrow L[i]$ 
             $i \leftarrow i + 1$ 
        else  $A[k] \leftarrow R[j]$ 
             $j \leftarrow j + 1$ 

```

Example: Bottom-up view for $n = 8$: [Heavy lines demarcate subarrays used in subproblems.]



Analyzing divide-and-conquer algorithms

Use a **recurrence equation** (more commonly, a **recurrence**) to describe the running time of a divide-and-conquer algorithm.

Let $T(n)$ = running time on a problem of size n .

- If the problem size is small enough (say, $n \leq c$ for some constant c), we have a base case. The brute-force solution takes constant time: $\Theta(1)$.
- Otherwise, suppose that we divide into a subproblems, each $1/b$ the size of the original. (In merge sort, $a = b = 2$.)
- Let the time to divide a size- n problem be $D(n)$.
- There are a subproblems to solve, each of size $n/b \Rightarrow$ each subproblem takes $T(n/b)$ time to solve \Rightarrow we spend $aT(n/b)$ time solving subproblems.
- Let the time to combine solutions be $C(n)$.
- We get the recurrence

$$T(n) = \begin{cases} \Theta(1) & \text{if } n \leq c, \\ aT(n/b) + D(n) + C(n) & \text{otherwise.} \end{cases}$$

Analyzing merge sort

For simplicity, assume that n is a power of 2 \Rightarrow each divide step yields two subproblems, both of size exactly $n/2$.

The base case occurs when $n = 1$.

When $n \geq 2$, time for merge sort steps:

Divide: Just compute q as the average of p and $r \Rightarrow D(n) = \Theta(1)$.

Conquer: Recursively solve 2 subproblems, each of size $n/2 \Rightarrow 2T(n/2)$.

Combine: MERGE on an n -element subarray takes $\Theta(n)$ time $\Rightarrow C(n) = \Theta(n)$.

Since $D(n) = \Theta(1)$ and $C(n) = \Theta(n)$, summed together they give a function that is linear in n : $\Theta(n) \Rightarrow$ recurrence for merge sort running time is

$$T(n) = \begin{cases} \Theta(1) & \text{if } n = 1, \\ 2T(n/2) + \Theta(n) & \text{if } n > 1. \end{cases}$$

Solving the merge-sort recurrence: By the master theorem we can show that this recurrence has the solution $T(n) = \Theta(n \lg n)$. [Reminder: $\lg n$ stands for $\log_2 n$.]

```
BinarySearch(A[0..n-1], value, low, high)
{
while (low <= high)
{
    mid = (low + high) / 2
    if (A[mid] > value)
        return BinarySearch(A, value, low, mid-1)
    else if (A[mid] < value)
        return BinarySearch(A, value, mid+1, high)
    else
        return mid          // found
}
return -1          // not found
}
```

Greedy Algorithms

Introduction

Similar to dynamic programming.

Used for optimization problems.

Idea: When we have a choice to make, make the one that looks best *right now*.
Make a *locally optimal choice* in hope of getting a *globally optimal solution*.

Greedy algorithms don't always yield an optimal solution. But sometimes they do. We'll see a problem for which they do. Then we'll look at some general characteristics of when greedy algorithms give optimal solutions.

Problem


- A town has a set of houses and a set of roads.
- A road connects 2 and only 2 houses.
- A road connecting houses u and v has a repair cost $w(u, v)$.
- **Goal:** Repair enough (and no more) roads such that
 1. everyone stays connected: can reach every house from all other houses, and
 2. total repair cost is minimum.

Model as a graph:

- Undirected graph $G = (V, E)$.
- **Weight** $w(u, v)$ on each edge $(u, v) \in E$.
- Find $T \subseteq E$ such that
 1. T connects all vertices (T is a **spanning tree**), and
 2. $w(T) = \sum_{(u,v) \in T} w(u, v)$ is minimized.

A spanning tree whose weight is minimum over all spanning trees is called a **minimum spanning tree**, or **MST**.

Example of such a graph [edges in MST are shaded] :



In this example, there is more than one MST. Replace edge (e, f) by (c, e) . Get a different spanning tree with the same weight.

Kruskal's algorithm


```
KRUSKAL( $V, E, w$ )  
   $A \leftarrow \emptyset$   
  for each vertex  $v \in V$   
    do MAKE-SET( $v$ )  
  sort  $E$  into nondecreasing order by weight  $w$   
  for each  $(u, v)$  taken from the sorted list  
    do if FIND-SET( $u$ )  $\neq$  FIND-SET( $v$ )  
      then  $A \leftarrow A \cup \{(u, v)\}$   
        UNION( $u, v$ )  
  return  $A$ 
```

Run through the above example to see how Kruskal's algorithm works on it:


(c, f) : safe
 (g, i) : safe
 (e, f) : safe
 (c, e) : reject
 (d, h) : safe
 (f, h) : safe
 (e, d) : reject
 (b, d) : safe
 (d, g) : safe
 (b, c) : reject
 (g, h) : reject
 (a, b) : safe

At this point, we have only one component, so all other edges will be rejected. [We could add a test to the main loop of KRUSKAL to stop once $|V| - 1$ edges have been added to A .]

Ex.



Solution:



Analysis

Initialize A : $O(1)$
First **for** loop: $|V|$ MAKE-SETs
Sort E : $O(E \lg E)$
Second **for** loop: $O(E)$ FIND-SETs and UNIONs

- Assuming the implementation of disjoint-set data structure,

$$O((V + E) \alpha(V)) + O(E \lg E) .$$

- Since G is connected, $|E| \geq |V| - 1 \Rightarrow O(E \alpha(V)) + O(E \lg E)$.
- $\alpha(|V|) = O(\lg V) = O(\lg E)$.
- Therefore, total time is $O(E \lg E)$.
- $|E| \leq |V|^2 \Rightarrow \lg |E| = O(2 \lg V) = O(\lg V)$.
- Therefore, $O(E \lg V)$ time. (If edges are already sorted, $O(E \alpha(V))$, which is almost linear.)

Dijkstra's algorithm

No negative-weight *edges*.

Essentially a weighted version of breadth-first search.

- Instead of a FIFO queue, uses a priority queue.
- Keys are shortest-path weights ($d[v]$).

Have two sets of vertices:

- S = vertices whose final shortest-path weights are determined,
- Q = priority queue = $V - S$.

DIJKSTRA(V, E, w, s)

INIT-SINGLE-SOURCE(V, s)

$S \leftarrow \emptyset$

$Q \leftarrow V$ \triangleright i.e., insert all vertices into Q

while $Q \neq \emptyset$

do $u \leftarrow \text{EXTRACT-MIN}(Q)$

$S \leftarrow S \cup \{u\}$

for each vertex $v \in \text{Adj}[u]$

do RELAX(u, v, w)

- $\pi[v] =$ predecessor of v on a shortest path from s .

All the shortest-paths algorithms start with INIT-SINGLE-SOURCE.

INIT-SINGLE-SOURCE(V, s)

for each $v \in V$

do $d[v] \leftarrow \infty$

$\pi[v] \leftarrow \text{NIL}$

$d[s] \leftarrow 0$

Relaxing an edge (u, v)

Can we improve the shortest-path estimate for v by going through u and taking (u, v) ?


RELAX(u, v, w)

if $d[v] > d[u] + w(u, v)$

then $d[v] \leftarrow d[u] + w(u, v)$

$\pi[v] \leftarrow u$

Example:



Order of adding to \mathcal{S} : s, y, z, x .

Analysis: If priority queue is implemented as binary heap : -

- Each Extract-Min operations then takes time $O(\lg V)$
- There are $|V|$ such operations
- The time to build the binary min-heap is $O(V)$
- Each Decrease-Key operation (implicit in Relax) takes time $O(\lg V)$, and there are at most $|E|$ such operations
- Total running time: $O((V+E) \lg V)$
 $\Rightarrow O(E \lg V)$

Because each vertex v is added to set S exactly once, each edge in the adjacency list $\text{Adj}[v]$ is examined in the for loop exactly once during the course of the algorithm. Since the total number of edges in all the adjacency lists is $|E|$, there are a total of $|E|$ iterations of this for loop. { Note: We are using aggregate analysis }

Scheduling

- Theorem: The greedy method always obtains an optimal solution to the job sequencing problem.

Algorithm for job scheduling using greedy strategy

{ Note: We want to find a feasible solution S whose profit $P(S)$ is as large as possible.
}

Sort the jobs in decreasing order of profits: $g_1 \geq \dots \geq g_n$

$d \leftarrow \max d_i$ // d_i stands for deadline for job i

for $t: 1..d$ // t stands for time slot

$S(t) \leftarrow 0$ // S stands for schedule

end for

for $i: 1..n$

Find the largest t such that ($S(t)=0$ and $t \leq d_i$), and let $S(t) \leftarrow i$

[Find the latest possible free slot meeting the deadline]

end for

{ Note: If $S(t) = 0$, then no job is scheduled at time t , g_i is a non-negative real no. representing the profit obtainable from job i . If $S(t)=i$, then job i is scheduled at time t , $1 \leq t \leq d_i$.}

Ex. Let $n=4$, $(P_1, P_2, P_3, P_4)=(100, 10, 15, 27)$ and $(d_1, d_2, d_3, d_4)=(2, 1, 2, 1)$

Solution:

Decreasing order of profits: 100 27 15 10

$d=2$

0	0
---	---

27	100
----	-----

Ex.: Let $n=5$, $(P_1, \dots, P_5)=(20, 15, 10, 5, 1)$ and $(d_1, \dots, d_5)=(2, 2, 1, 3, 3)$. Find the optimal solution.

Solution: The optimal solution is :

15	20	5
----	----	---

Ex. Given below is a problem of scheduling unit-time tasks with deadlines and penalties for a single processor.

Tasks

i	1	2	3	4	5	6	7
d_i	4	2	4	3	1	4	6
w_i	70	60	50	40	30	20	10

Solution:

40	60	50	70		10
----	----	----	----	--	----

Total penalty incurred is $w_5 + w_6 = 50$